A Century and a Half at the Ballot Box

Regimes, linkage and data gaps in Ventura County election returns, 1876-2026

Nikhil Chandrashaker | California Lutheran University, Data Science | October 23, 2026

Abstract

This study turns a bundle of Ventura County, California election files into a research database and asks three questions. First, how stable has the county's partisan lean been? Second, how closely do gubernatorial results follow the preceding presidential vote? Third, how reliable is the record itself? Across 38 presidential elections (1876-2024) a change-point model finds 5 regimes, and the winning party changed 9 times. The two-party margin has moved from a long Republican era, through a New Deal Democratic era (1932-1964) and a Republican era (1968-1988), to a Democratic era since 1992 (mean margin D+8.2 points). Governor results (n = 16) are only loosely tied to the previous presidential result (R^2 = 0.19, leave-one-out error 8.28 points), so a simple linear carry-over model is not a usable forecast. The audit of the supplied files found 16 logged data-quality issues, including a mislabelled 1912 winner and a probable candidate-name error in the 1962 governor race. Only presidential and gubernatorial elections are deep; local records exist for 2020-2026 only. The paper documents the models, the schema that keeps raw results, source notes and missing-data records together, and a roadmap for adding older local elections.

1. Introduction

County election returns are a rare long record of how a place votes, but they are usually scattered across spreadsheets, scanned statements of vote and web tables. Ventura County was established on January 1, 1873. Its first presidential returns in the supplied data are from 1876, when it was a small agricultural county of about 1,200 voters; by 2020 the presidential vote had grown to more than 420,000. This study builds an auditable database from the supplied files and uses it for three jobs: an interactive election-history dashboard, a candidate and election explorer, and a research record that keeps every unverified or missing item visible.

Contributions: (1) a normalised relational schema with explicit review status per contest; (2) a documented audit of the source files; (3) a regime analysis of the presidential margin; (4) a test of how much the previous presidential vote explains the next governor race; and (5) a plan for extending the database to city, school-district and special-district contests.

2. Data and provenance

The bundle has three views of the same material. The *master* file holds 366 candidate-level rows with only a year for each election date. The *cleaned* file holds the same 366 rows with full dates, a stage (primary or general), a quality flag and a source label. The *analytics* file holds 154 yearly rows (1873-2026) with presidential and gubernatorial totals, swings, competitiveness bands and one voter-registration snapshot (May 18, 2026). All analysis here starts from the cleaned file; the analytics file is used for cross-checks.

The cleaned file contains 92 contests across 23 offices, 295 named candidates and 15 party labels. Presidential general elections account for 117 rows and gubernatorial rows for 98. Rows come from four labelled sources: an earlier presidential source (117 rows), an earlier gubernatorial source (32), an earlier county and local extract (45), and a California Secretary of State extract of the June 2, 2026 primary Statement of Vote (172). The

bundle supplied no URLs, so the sources table records each source as unverified with an empty link field. Nothing here should be cited as verified against official statements of vote.

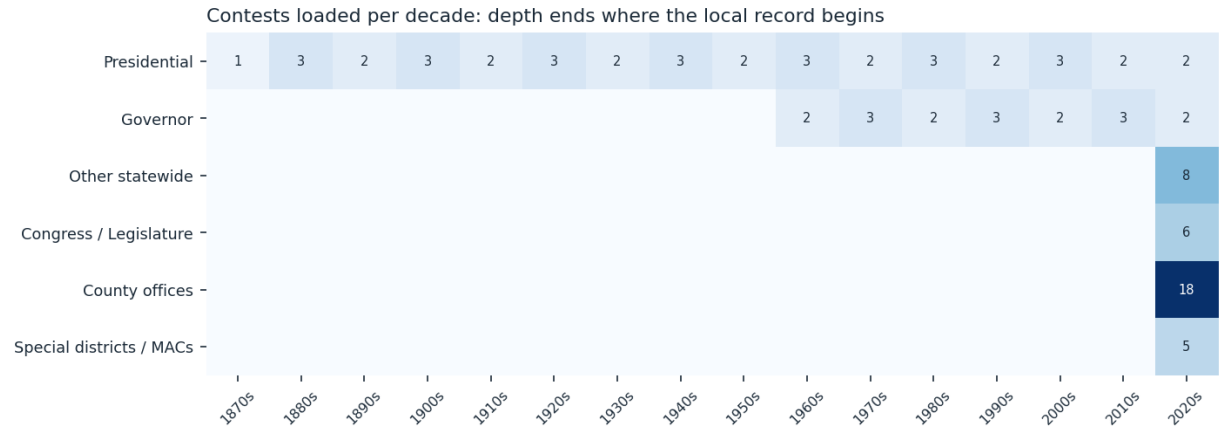


Figure 1. Contests loaded per decade. Presidential and gubernatorial depth contrasts with a local record that starts in 2020.

3. Research database design

The schema separates what was observed from what is known about its quality. *contests* holds one row per contest with a review status; *results* holds candidate rows; *candidates* holds normalised names (the pooled “Other” rows are not treated as a person); *sources* holds provenance with a verified flag; *data_quality_issues* and *missing_data* are first-class tables rather than footnotes. Views provide winners, candidate histories and coverage. Values are never overwritten: an error is logged beside the value it concerns.

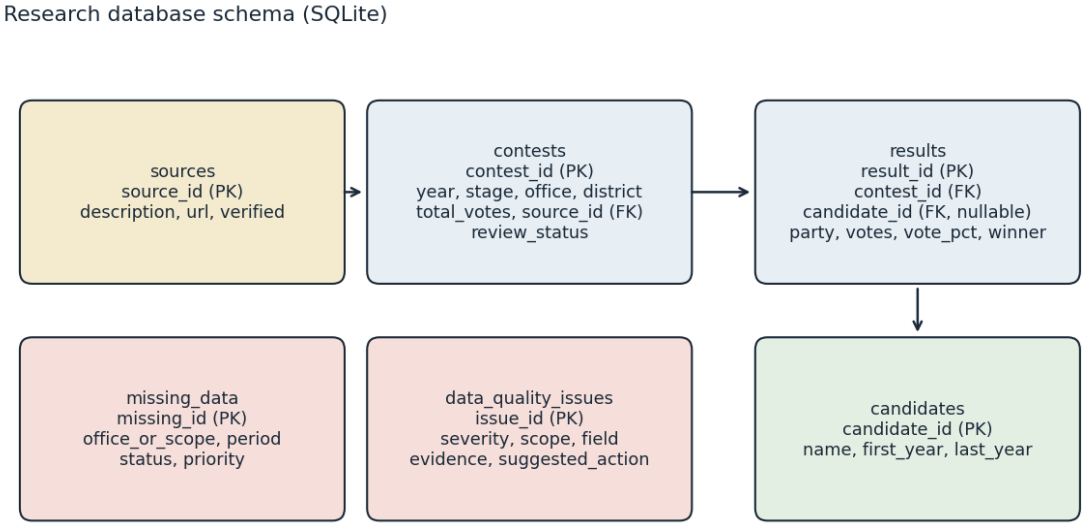


Figure 2. Research database schema (SQLite).

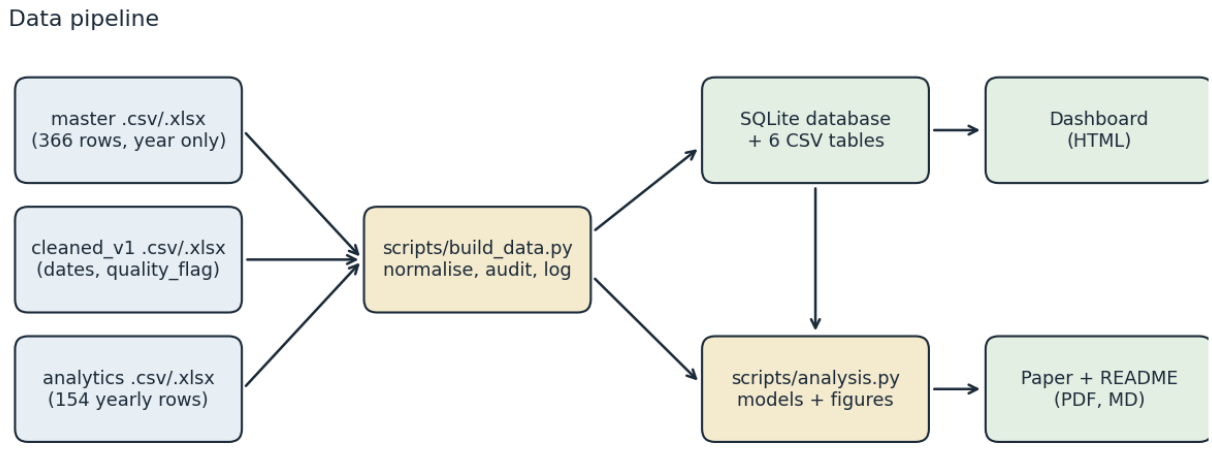


Figure 3. Pipeline from the three supplied files to database, models, dashboard and paper.

4. Methods

Two-party margin. For each presidential general election the margin is 100 x (D - R) / (D + R) in percentage points, where D and R are Democratic and Republican votes. In 1912 the Progressive candidate Theodore Roosevelt, a Republican breakaway who took 54.9% in the county, is added to the Republican side so that the series measures Democratic versus Republican lineage. Third-party votes (Perot in 1992, La Follette in 1924) are excluded from the denominator, so margins in those years overstate the two-party gap.

Regimes. The margin series is split into k contiguous segments with segment means minimising the sum of squared errors, solved exactly by dynamic programming; k is chosen by a BIC criterion (n log(RSS / n) + 2k log n). This describes eras without assuming a smooth trend.

Linkage. For each gubernatorial general election (1962-2022, n = 16) the Democratic two-party share is regressed on the Democratic two-party share of the presidential election two years earlier. Fit is summarised by R^2 and leave-one-out RMSE computed from hat-matrix residuals.

Midterm gap. The gap is gubernatorial share minus previous presidential share. A second regression adds an indicator for the sitting president's party (+1 Democratic, -1 Republican). A midterm penalty would appear as a negative coefficient.

Turnout and volatility. Log presidential votes are regressed on year from 1932 to 2024. Volatility is the mean absolute election-to-election swing in margin by era.

5. Results

5.1 Five regimes in the presidential margin

Period	Elections	Mean margin (pp)	SD (pp)	Reading
1876-1916	11	-13.8	10.8	Republican, with a 1912 Progressive split
1920-1928	3	-58.0	15.5	Republican landslides
1932-1964	9	+12.2	11.4	Democratic, with Republican wins in 1952 and 1956

Period	Elections	Mean margin (pp)	SD (pp)	Reading
1968-1988	6	-25.0	12.4	Republican
1992-2024	9	+8.2	9.3	Democratic, narrow before 2008

The county's winning party changed 9 times across the 38 elections (Republican 23, Democratic 14, Progressive 1). The most Democratic result is 1936 (D+27.7); the most Republican is 1924 (R-72.5, inflated by La Follette's exclusion). Eight elections were decided by under five points (1876, 1896, 1956, 1960, 1992, 1996, 2000, 2004); four of them fall in 1992-2004, the narrow stretch before the county's Democratic margin widened in 2008.

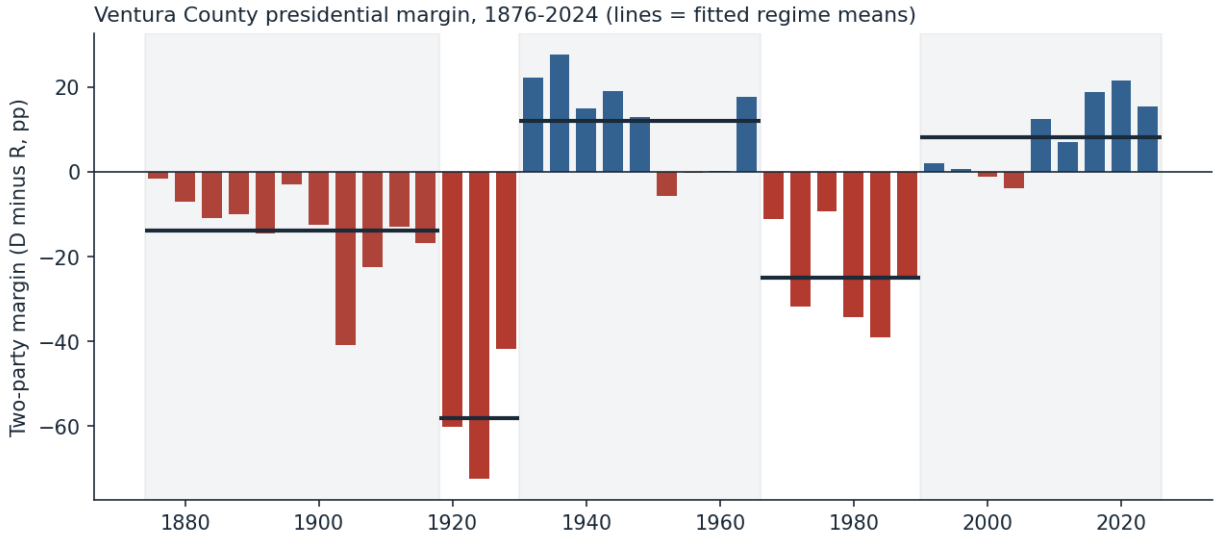


Figure 4. Presidential two-party margin with fitted regime means. Blue bars are Democratic pluralities, red are Republican.

Mean absolute swing between consecutive elections was 14.0 points in 1876-1928, 16.6 in 1932-1956, 16.7 in 1960-1988 and 8.3 in 1992-2024. The modern era is the least volatile even though it contains the county's closest contests, which suggests a stable near-even electorate that gave way to a steadier Democratic lean after 2004.

5.2 Governor results and the previous presidential vote

The linear model gives Democratic governor share = 27.94 + 0.38 x previous presidential Democratic share (slope SE 0.209, n = 16). It explains 19% of variance, and the leave-one-out RMSE is 8.28 points, about as large as the post-1992 regime's average presidential margin (D+8.2). A slope near 0.4 rather than 1 means gubernatorial results do not simply copy presidential partisanship: candidate quality, midterm dynamics and 1966-1994 Republican gubernatorial strength all intervene.

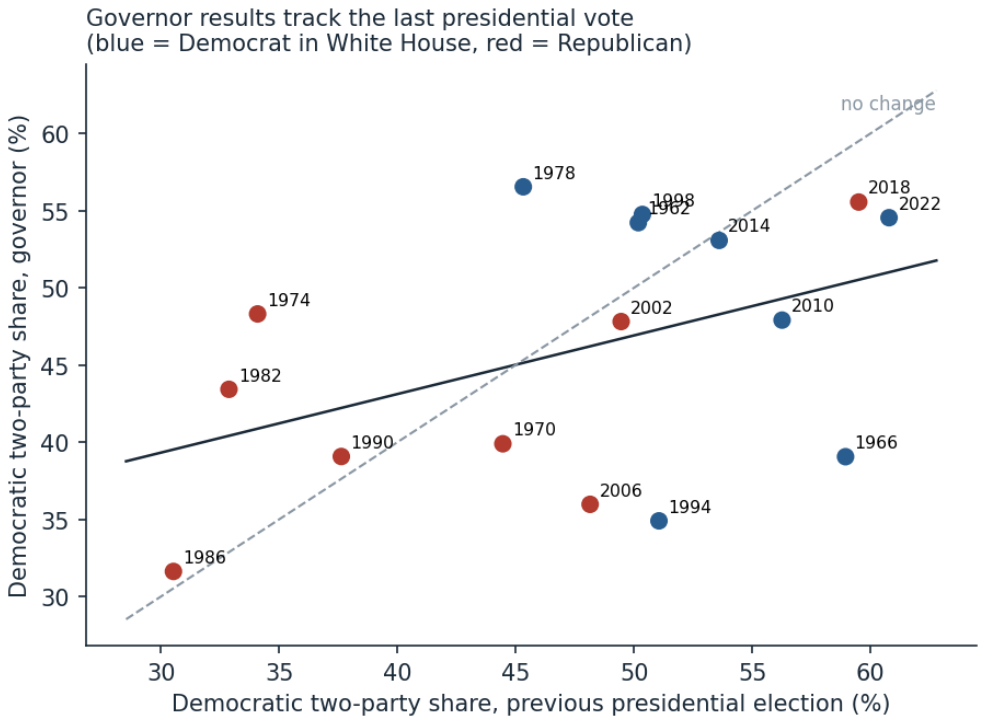


Figure 5. Governor Democratic share against the previous presidential Democratic share. Points above the dashed line are Democratic gains relative to the last presidential vote.

5.3 Midterm gap

Raw gaps are larger when a Democrat is president: the mean gap is -3.9 points under Democratic presidents and +0.6 under Republican presidents. A regression that also controls for the previous presidential share finds a coefficient of +1.89 points on the sitting president's party (SE 2.48, p = 0.461), which is small, has the opposite sign to a midterm penalty and is not statistically distinguishable from zero. The raw pattern and the controlled estimate disagree, so with 16 observations the evidence for a midterm penalty is inconclusive.

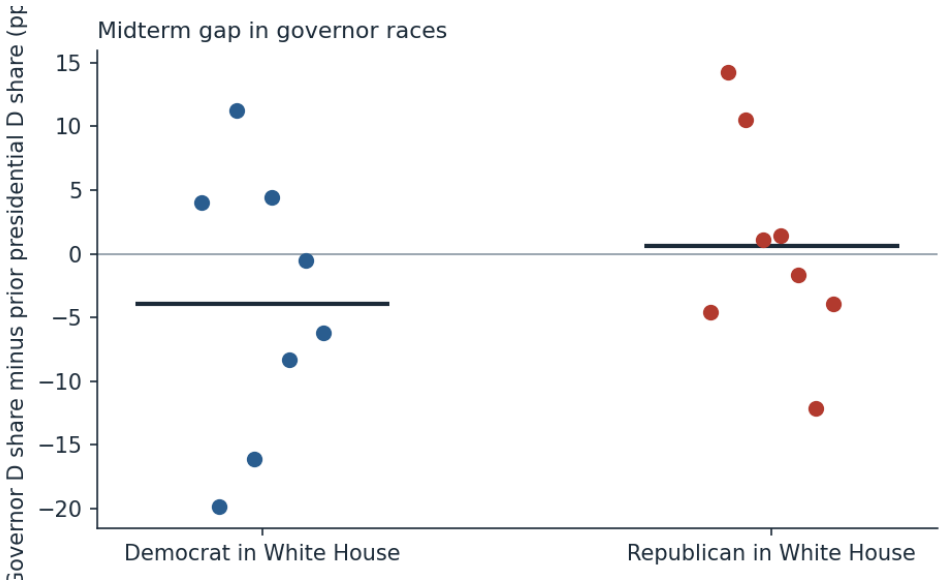


Figure 6. Gap between governor and previous presidential Democratic share, by party of the sitting president.

5.4 Turnout

From 1932 to 2024 presidential turnout grew about 42.5% per decade (log-linear fit, $R^2 = 0.933$). Votes peaked in 2020 at 422,825 and fell to 387,595 in 2024, a decline of about 8%. No registration series exists in the bundle, so participation rates cannot be computed.

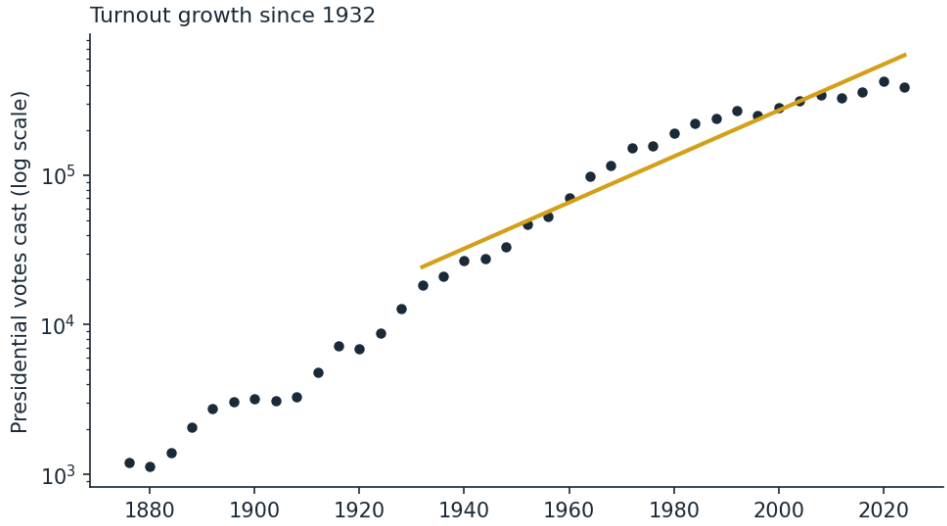


Figure 7. Presidential votes cast on a log scale with fitted growth.

6. A cautionary baseline for the 2026 governor race

Applying the linkage model to the 2024 presidential result (Democratic two-party share 57.8%) with a Republican president gives a predicted Democratic share of 46.8% with a 95% prediction interval of 26.5% to 67.0%. That is not a forecast. The interval spans nearly the whole plausible range, and the point estimate sits below every gubernatorial result since 2014 (53.1%, 55.6% and 54.5% Democratic two-party share in 2014,

2018 and 2022). The model is stationary while the county is not: it is trained mostly on a Republican gubernatorial era. The exercise shows why regime-aware or candidate-level models, and not a single line fitted across sixty years, are needed. The November 3, 2026 general election is not in the data.

7. Data-quality audit

Every issue is stored in *data_quality_issues* with evidence and a suggested action. The most consequential are:

ID	Severity	Scope	Finding
DQ01	high	President 1912	Analytics file labels the 1912 winner 'Democratic' with a +42.1 pp margin, but the cleaned file shows Theodore Roosevelt (Progressive) first with 54.9% to Wilson's 43.6%.
DQ02	high	Governor 1962	Democratic nominee listed as Glenn M. Anderson (Democratic), but Anderson was the sitting Lieutenant Governor; the Democratic nominee for governor in 1962 was Pat Brown. Vote totals may belong to a different contest.
DQ03	medium	President 2024	Cleaned file reports 387,595 total votes; analytics file reports 387,704 (109 more) with identical Democratic and Republican counts.
DQ04	medium	Supervisor D1 2020 primary	Rows flagged INCOMPLETE: only two candidates supplied and they do not reconcile to the reported contest total.
DQ05	medium	Local contests with >1 winner	5 contests have several winners (multi-seat boards/districts) but the number of seats is not recorded, so winner flags cannot be validated.
DQ06	medium	Supervisor D5 2020 primary	No candidate flagged as winner. No one cleared 50% (leader Flynn 29.2%), consistent with a runoff, but runoff logic is not encoded.
DQ07	medium	All primaries	In primaries 'winner' means top finisher, not nominee or office-holder (for example, top-two and party primaries).
DQ08	medium	All rows	Incumbent column is empty in all 366 rows.

The 1912 and 1962 findings matter most. In 1912 the analytics file reports a Democratic winner with a 42-point margin, while the cleaned file shows Roosevelt with 54.9%; using the analytics label would have created a false Democratic county in 1912. In 1962 the Democratic candidate is named Glenn M. Anderson, who was the sitting lieutenant governor, whereas the Democratic gubernatorial nominee that year was Pat Brown; the vote totals are retained but the name should be checked against the county statement of vote. I flag this from historical knowledge, not from a source document in the bundle.

8. Limitations

Sixteen gubernatorial observations support only coarse inference. Two-party shares ignore third parties and top-two primary dynamics. The 1912 grouping is a modelling choice. All values are carried over from supplied files that are not verified against official statements of vote. Winner flags in primaries mean top finisher, and the incumbent field is empty. No precinct, city or turnout-rate data exist, so the analysis cannot separate changes in who votes from changes in how they vote.

9. Roadmap for older and local elections

Priority order, from the missing-data table: (1) gubernatorial generals 1879-1958, about 21 contests whose years are inferred from the four-year cycle and need verification; (2) general-election results for congressional, legislative and statewide offices before 2026; (3) Board of Supervisors and other county offices before 2020; (4) city councils and mayors; (5) school-district and community-college boards; (6) special districts; (7) ballot measures, registration history and turnout. For each added contest: record the source URL and retrieval date, set

the review status, add the seat count for multi-seat races, and fill the incumbent field. A contest moves to verified only when its candidate rows reconcile to the official contest total.

Data and code

Supplied files are in data/raw. Scripts build_data.py, analysis.py, build_dashboard.py and make_paper.py regenerate every table, figure, the dashboard and this PDF. Source names are as supplied: Wikimedia Commons Ventura County presidential returns; a Ventura County gubernatorial results table; California Secretary of State Statement of Vote for the June 2, 2026 primary; California Secretary of State Report of Registration, May 18, 2026.